How I used AI on this challenge

The brief invited any tools, including AI. Here is exactly what I did, what AI did, and how I checked it.

Time: about **4 hours** of my own time over two sessions. Part 3 is the n8n workflow I built on 13 October for my first recruiter call, before this exercise, and the narrated video I made of it. I didn't change it for this case.

What came from me

These are the decisions the work is built on. I specified them, and corrected the AI where its first version didn't match how I run renewals. I also used it as a sparring partner to check whether I was missing an option.

Area	My decision
Forecast	Renewal Status × days to renewal drives the probability. Later statuses rise as the date approaches, early ones fall. Statuses are set from what we know today, with a second view after this week's actions ("forecast of the forecast").
Overdue renewals	A separate league: they only count if they close within the quarter. Otherwise they count as zero for that quarter.
Auto-renewals	Their own statuses: blockers, at risk before or after the notice deadline. Over-usage opens an upsell opportunity, never an automatic charge. Under-usage flags the CSM.
Product lines and churn	Each line is renewal, upsell, downgrade or churn. Original ARR is kept next to new ARR, and Salesforce recalculates everything else. Churn sits on the opportunity: **regrettable** means revenue lost for good; **non-regrettable** means revenue kept elsewhere at n8n (e.g. self-hosted → Cloud). Upsell is its own opportunity.
Commercial rules	The 5% uplift is already in the renewal value (never added on top), announced in one bulk notice each January. The price is locked from year 2 only for upfront multi-year deals. The Renewals Manager owns the numbers. A new renewal opportunity is created when the current one closes won.
Brightloop	Two structures instead of a discount: downsize + multi-year lock, or keep the plan with a small buffer. Net 60 only with upfront payment; top SLA tier.
Automation	My own workflow: amounts come from the source of truth, never from the model; a human approves every customer-facing email; risky accounts get an alert to a person, not an email.

What AI (Claude) did

- **Built from my specifications:** the Excel tracker (formulas, tabs, formatting) and the PDFs.
- **Drafted** text I then edited: the email to Sam, the playbook, the write-ups.
- **Did and checked the maths:** the policy C approval thresholds and the multi-year discount maths, e.g. why a 10% + 5% ramp is 7.4% overall, not 15%.
- **Researched** facts about n8n (self-hosted and Cloud editions, pricing by executions) and flagged the assumptions it couldn't verify. One example: the price of a smaller Enterprise tier.

How I checked it

- I reviewed every number in the tracker and documents. The totals tie out across files.
- **Where AI got it wrong, and I corrected it:**
  - its first forecast relied on judgement instead of my status × countdown model;
  - it put churn type on the product lines instead of the opportunity;
  - it proposed waiving the uplift on 2-year deals without upfront payment;
- **An AI failure I caught while testing an extension of the workflow for this case:** an AI-drafted email leaked internal notes ("budget cut", "usage down 10%"). That's why every customer-facing draft goes to a person for approval.

My principle

The same as in the AI agents I build at work: numbers come from the tool of record, never from the model; AI is used where language helps; a human approves anything customer-facing; and autonomy is earned with a track record. I know what was built, why and how, because I designed it. AI let me build it in hours instead of days.